

# Using Commit Hooks

## Contents

1. Overview
  2. Configuration
  3. Pre-Commit Hooks
  4. Post-Commit Hooks
- 

## Overview

Pre- and post-commit hooks are invoked before or after a `riak_object` is persisted and can greatly enhance the functionality of any application. Commit hooks can:

- allow a write to occur with an unmodified object
- modify the object
- Fail the update and prevent any modifications

Post-commit hooks are notified after the fact and should not modify the `riak_object`. Updating `riak_objects` in post-commit hooks can cause nasty feedback loops which will wedge the hook into an infinite cycle unless the hook functions are carefully written to detect and short-circuit such cycles.

Pre- and post-commit hooks are defined on a per-bucket basis and are stored in the target bucket's properties. They are run once per successful response to the client.

## Configuration

Configuring either pre- or post-commit hooks is very easy. Simply add a reference to your hook function to the list of functions stored in the correct bucket property. Pre-commit hooks are stored under the bucket property *precommit*. Post-commit

hooks use the bucket property *postcommit*.

Pre-commit hooks can be implemented as named Javascript functions or as Erlang functions. The configuration for each is given below:

Javascript: `{"name": "Foo.beforeWrite"}` Erlang: `{"mod": "foo", "fun": "beforeWrite"}`

Post-commit hooks can be implemented in Erlang only, and is described in more detail under [Advanced Commit Hooks](#). The reason for this restriction is Javascript cannot call Erlang code and, thus, is prevented from doing anything useful. This restriction will be revisited when the state of Erlang/Javascript integration is improved. Post-commit hooks use the same function reference syntax as pre-commit hooks.

See [Advanced MapReduce](#) for steps to define your own pre-defined Javascript named functions.

# Pre-Commit Hooks

## API & Behavior

Pre-commit hook functions should take a single argument, the `riak_object` being modified. Remember that deletes are also considered “writes” so pre-commit hooks will be fired when a delete occurs. Hook functions will need to inspect the object for the *X-Riak-Deleted* metadata entry to determine when a delete is occurring.

Erlang pre-commit functions are allowed three possible return values:

- A `riak_object` – This can either be the same object passed to the function or an updated version. This allows hooks to modify the object before they are written.
- `fail` – The atom *fail* will cause Riak to fail the write and send a 403 Forbidden along with a generic error message about why the write was blocked.
- `{fail, Reason}` – The tuple `{fail, Reason}` will cause the same

behavior as in #2 with the addition of `Reason` used as the error text.

Errors that occur when processing Erlang pre-commit hooks will be reported in the `sasl-error.log` file with lines that start with “problem invoking hook”.

#### Erlang Pre-commit Example

##### Erlang

```
%% Limits object values to 5MB or smaller
precommit_limit_size(Object) ->
    case erlang:byte_size(riak_object:get_value(Object)) of
        Size when Size > 5242880 -> {fail, "Object is larger than 5MB"}
        _ -> Object
    end.
```

Javascript pre-commit functions should also take a single argument, the JSON encoded version of the `riak_object` being modified. The JSON format is exactly the same as Riak's MapReduce. Javascript pre-commit functions are allowed three possible return values:

- A JSON encoded Riak object – Aside from using JSON, this is exactly the same as #1 for Erlang functions. Riak will automatically convert it back to its native format before writing.
- `fail` – The Javascript string “fail” will cause Riak to fail the write in exactly the same way as #2 for Erlang functions.
- `{"fail": Reason}` – The JSON hash will have the same effect as #3 for Erlang functions. Reason must be a Javascript string.

#### Javascript Pre-commit Example

##### Javascript

```
// Makes sure the object has JSON contents
function precommitMustBeJSON(object){
    try {
        Riak.mapValuesJson(object);
        return object;
    } catch(e) {
        return {"fail": "Object is not JSON"};
    }
}
```

# Chaining

The default value of the bucket *precommit* property is an empty list. Adding one or more pre-commit hook functions, as documented above, to the list will cause Riak to start evaluating those hook functions when bucket entries are created, updated, or deleted. Riak stops evaluating pre-commit hooks when a hook function fails the commit.

## Example

Pre-commit hooks can be used in many ways in Riak. One such way to use pre-commit hooks is to validate data before it is written to Riak. Below is an example that uses Javascript to validate a JSON object before it is written to Riak.

### Javascript

```
//Sample Object
{
  "user_info": {
    "name": "Mark Phillips",
    "age": "25",
  },
  "session_info": {
    "id": 3254425,
    "items": [29, 37, 34]
  }
}

var PreCommit = {
  validate: function(obj){

    // A delete is a type of put in Riak so check and see if the
    // operation is doing

    if (obj.values[[0]][[ 'metadata' ]][[ 'X-Riak-Deleted' ]])
      return obj;
  }
}
```

```

    // Make sure the data is valid JSON
    try{
        data = JSON.parse(obj.values[[0]].data);
        validateData(data);

    }catch(error){
        return {"fail": "Invalid Object: "+error}
    }
    return obj;
}

function validateData(data){
    // Validates that user_info object is in the data
    // and that name and age aren't empty, finally
    // the session_info items array is checked and validated
    // being populated

    if(
        data.user_info != null &&
        data.user_info.name != null &&
        data.user_info.age != null &&
        data.session_info.items.length > 0
    ){
        return true;
    }else{
        throw( "Invalid data" );
    }
}

```

# Post-Commit Hooks

## API & Behavior

Post-commit hooks are run after the write has completed successfully. Specifically,

the hook function is called by `riak_kv_put_fsm` immediately before the calling process is notified of the successful write. Hook functions must accept a single argument, the `riak_object` instance just written. The return value of the function is ignored. As with pre-commit hooks, deletes are considered writes so post-commit hook functions will need to inspect object metadata for the presence of *X-Riak-Deleted* to determine when a delete has occurred. Errors that occur when processing post-commit hooks will be reported in the `sasl-error.log` file with lines that start with “problem invoking hook”.

Example

Erlang

```
%% Creates a naive secondary index on the email field of a
postcommit_index_on_email(Object) ->
    %% Determine the target bucket name
    Bucket = erlang:iolist_to_binary([riak_object:bucket(Object),
    %% Decode the JSON body of the object
    {struct, Properties} = mochijson2:decode(riak_object:get(Object)),
    %% Extract the email field
    {<<"email">>, Key} = lists:keyfind(<<"email">>, 1, Properties),
    %% Create a new object for the target bucket
    %% NOTE: This doesn't handle the case where the
    %%       index object already exists!
    IndexObj = riak_object:new(Bucket, Key, <<>>, %% no object exists
                               dict:from_list(
                                   [
                                       {<<"content-type">>, "text/plain"},
                                       {<<"Links">>,
                                           [
                                               {{riak_object:bucket(Object),
                                               ])),
    %% Get a riak client
    {ok, C} = riak:local_client(),
    %% Store the object
    C:put(IndexObj).
```

## Chaining

The default value of the bucket *postcommit* property is an empty list. Adding one or more post-commit hook functions, as documented above, to the list will cause Riak to start evaluating those hook functions immediately after data has been created, updated, or deleted. Each post-commit hook function runs in a separate process so it's possible for several hook functions, triggered by the same update, to execute in parallel. *All post-commit hook functions are executed for each create, update, or delete.*

## These May Also Interest You

- [Five-Minute Install](#)
- [Advanced Secondary Indexes](#)
- [Replication Properties](#)
- [Advanced Commit Hooks](#)
- [Advanced MapReduce](#)
- [Advanced Search](#)